


CZECH TECHNICAL UNIVERSITY IN PRAGUE

FACULTY OF ELECTRICAL ENGINEERING
DEPARTMENT OF CYBERNETICS



Classification of journalistic quality of articles

Bachelor's Thesis

Olga Pimenova

Prague, May 2025

Study programme: Open Informatics
Branch of study: Artificial Intelligence and Computer Science

Supervisor: Ing. Jan Šedivý, CSc.

Acknowledgments

I sincerely thank my academic supervisor, Ing. Jan Šedivý, CSc., for his expert guidance, patience, and encouragement. His profound knowledge of advanced technologies and dedication to academic excellence significantly shaped the direction and quality of my research.

I am equally grateful to my workplace supervisor, Radek Tomáš, for his practical insights and mentorship during the implementation phase. His feedback effectively bridged theoretical concepts with real-world applications, enhancing the project's outcomes.

Special thanks to the Seznam.cz recommendation systems team for providing technical resources, collaborative opportunities, and industry perspectives that strengthened the experimental aspects of this thesis.

Finally, I owe immense gratitude to my family for their unwavering emotional support, sacrifices, and belief in me. Their encouragement during challenging moments was crucial to completing this journey.

Swap this for the Thesis Assignment, when you have it from the department.

Declaration

I declare that this thesis was independently developed and that all sources of information used are cited in accordance with the university's ethical guidelines for thesis preparation. Additionally, I disclose that I utilized the AI tool ChatGPT to assist in translating certain sentences from my native language into English and to rephrase them in a style more appropriate for academic writing.

Prague, May 2025

Abstract

TODO:

Keywords Unmanned Aerial Vehicles, Automatic Control

Abstrakt

TODO:

Klíčová slova Bezpilotní Prostředky, Automatické Řízení

Abbreviations

GPS Global Positioning System

LiDAR Light Detection and Ranging

UAV Unmanned Aerial Vehicle

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Problem Formulation	2
1.3	Structure of the Thesis	2
1.4	Related Work	2
1.4.1	Quality Assessment Frameworks in Journalism	3
1.4.2	Computational Approaches to Text Quality Assessment	3
1.4.3	Machine Learning for Classification of Media Content	3
1.4.4	Multi-class Classification in Natural Language Processing	3
1.4.5	AI in Journalism and Quality Implications	4
1.4.6	Linguistic Features for Quality Assessment	4
1.4.7	Research Gaps	4
1.5	Contributions	4
2	How to write thesis in LaTeX	5
2.1	Versioning with git	5
2.2	Forming paragraphs	5
2.3	Linguistic anti patterns	5
2.3.1	Narrative	5
2.3.2	Pronouns	5
2.4	Mathematical notation with LaTeX	5
2.4.1	Common errors	6
2.4.2	Equations	6
2.5	Using footnotes	7
2.6	Referencing document elements	7
2.7	Abbreviations with Acronym	7
2.8	Units of measurements with Siunitx	7
2.9	Hyphens and dashes	7
2.10	Double quotation marks	8
2.11	2D Diagrams with Tikz	8
2.12	Data plots with PGFPlots	8
2.13	3D Plots with Sketch	9
2.14	Image collages with Subfig	9
2.15	Citations with Biblatex	9
2.16	Image overlays with Tikz	9
2.17	General tips	10
3	Conclusion	12

4	References	13
A	Appendix A	14

1 Introduction

1.1 Motivation

The digital content revolution has transformed journalism into a high-velocity industry where platforms must simultaneously manage exponential content growth and maintain rigorous quality standards. Seznam.cz, serving over 5 million daily users, currently relies on human editors to manually assess only 20% of submitted articles - an unsustainable practice given the platform's 50% annual content growth rate. This manual process introduces significant challenges:

- **Scalability limitations:** With article submissions increasing by 50% annually, human editors cannot maintain pace without compromising review depth
- **Consistency issues:** Inter-editor agreement analysis reveals 20% variance in quality ratings for identical content
- **Quality degradation risks:** Unreviewed articles risk platform credibility through potential misinformation or poorly structured content

These limitations highlight the need for more efficient and consistent methods for evaluating journalistic quality. Automated quality classification addresses these challenges by providing:

- Real-time assessment for 100% of submissions
- Consistent application of editorial guidelines
- Significantly reduce the workload on human editors
- 24/7 operational capacity without human fatigue
- Data-driven insights for continuous quality improvement

Therefore, the primary motivation for this research is to develop an automated system capable of classifying the journalistic quality of articles, mimicking the complex decision-making process of experienced human editors.




Figure 1.1: Main page Seznam.cz




Figure 1.2: Current Editorial Workflow

■ 1.2 Problem Formulation

The core problem addressed in this thesis is the multi-class supervised classification of journalistic articles based on their quality. Our goal is to develop a computational model that automatically assigns an article to one of four predefined quality categories, reflecting the standards used by editors at Seznam.cz.

- Class 0 (Insufficient): Low-quality content that lacks structure and depth.
- Class 1 (Short but Sufficient): Lacks depth and analysis, often recycling known information with limited sources.
- Class 2 (Standart): Reliable news articles with at least five paragraphs and two references.
- Class 3 (High-Quality): Long, detailed, well-structured articles with visuals, subheadings, and references.

In the context of this thesis, **”journalistic quality”** is defined by a combination of factors that human editors consider. These include textual properties like readability, complexity, and style; structural attributes such as organization, paragraphing, and the use of headings; content characteristics like depth of analysis, factual accuracy, and source referencing; and metadata signals like the publishing channel. The aim is to create a model that learns to identify the patterns associated with these factors to predict the appropriate quality category. A significant challenge within this problem is the inherent imbalance in real-world data, where lower-quality categories are often much more frequent than the highest-quality ones.

■ 1.3 Structure of the Thesis

■ 1.4 Related Work

The automated classification of journalistic quality is a multidisciplinary challenge that integrates natural language processing (NLP), machine learning (ML), and media studies. This section critically reviews existing literature across these fields, highlighting significant

contributions and identifying gaps that this thesis addresses, particularly in the context of the Czech language media environment and platforms like Seznam.cz.

■ Quality Assessment Frameworks in Journalism

Traditional journalistic quality assessment relies on human evaluation guided by professional standards. **lacy2015** established eight core dimensions of journalistic quality, including accuracy, context, and sourcing diversity. While comprehensive, their framework was designed for manual evaluation, limiting its applicability to automated systems. Recent work by **calvo2024** extends these principles to AI-assisted journalism, emphasizing transparency, accountability, and ethical considerations as critical criteria for maintaining quality in digital news production. Their study highlights that AI can enhance efficiency but risks reducing nuance and context, necessitating robust evaluation frameworks.

■ Computational Approaches to Text Quality Assessment

Automated text quality assessment provides foundational insights for journalistic applications. **pitler2008** demonstrated correlations between syntactic features and human perceptions of writing quality, using metrics like sentence complexity and coherence. However, their approach was general and did not address the specific stylistic and ethical requirements of journalistic texts. More recent advancements, as noted in a 2024 article by **quantum2024**, leverage NLP techniques such as grammar analysis, sentiment analysis, and text cohesion metrics to evaluate content quality automatically. These methods offer promising tools for journalism but require adaptation to domain-specific standards.

■ Machine Learning for Classification of Media Content

Machine learning has significantly advanced media content classification. **potthast2018** developed methods to detect style-based bias in news articles, focusing on political orientation. Their work illustrates the potential of ML to identify subtle textual patterns but does not directly address overall quality. Similarly, **gonzalez2021** applied BERT-based architectures to classify news articles by credibility, achieving accuracy rates above 85%. However, their binary classification approach (credible vs. non-credible) does not capture the nuanced spectrum of journalistic quality, which often requires multi-class models.

■ Multi-class Classification in Natural Language Processing

Multi-class classification in natural language processing (NLP) is a critical technique for categorizing text into multiple predefined classes, making it highly relevant for assessing the quality of journalistic content across various dimensions, such as accuracy, depth, and balance. A comprehensive review by **minaee2020** examines over 150 deep learning models applied to text classification tasks, including sentiment analysis, news categorization, and question answering. The authors highlight the effectiveness of transformer-based architectures, such as BERT and its variants, which leverage large-scale pretraining and fine-tuning to achieve state-of-the-art performance in complex multi-class classification problems. These models are particularly promising for evaluating journalistic quality, as they can capture semantic nuances and contextual relationships within news articles. The review also summarizes over 40 popular datasets, some of which focus on news categorization, providing valuable benchmarks for developing quality assessment systems.

Similarly, **schmidt2019** introduced a multi-class approach for classifying Wikipedia article quality using document embeddings and content features. Their method, which achieved high accuracy in distinguishing quality levels, demonstrates the potential of semantic representations for quality assessment, adaptable to journalistic contexts.

■ AI in Journalism and Quality Implications

The integration of AI in journalism has transformed news production, as evidenced by recent studies. **newman2025** report that 87% of newsrooms have adopted generative AI for tasks like content creation (77% deem it important) and automation. However, they note concerns about "AI slop" (low-quality AI-generated content) and the need for transparency, such as adopting C2PA metadata standards. A 2025 mini-review by **frontiers2025** found that AI has revolutionized news production but poses ethical challenges, including reduced nuance in storytelling. These findings underscore the need for automated quality assessment tools that align with journalistic values.

■ Linguistic Features for Quality Assessment

Identifying linguistic features that correlate with quality is a key research area. **schmidt2019** utilized document embeddings to capture semantic content, combined with features like section length and readability scores, to assess article quality. In journalism, similar features—such as Flesch-Kincaid readability, lexical diversity, and syntactic complexity—can indicate quality but must be tailored to journalistic standards. The Czech language, with its rich morphology and inflectional system, presents unique challenges for NLP tools, as standard models designed for English may underperform.

■ Research Gaps

Despite progress, several gaps persist in the literature. First, most studies focus on English-language content, with limited exploration of languages like Czech, where morphological complexity affects NLP performance. Second, many approaches rely on binary classification, which oversimplifies the multi-dimensional nature of journalistic quality. Third, few studies address the specific needs of news aggregation platforms like Seznam.cz, where high content volume and diversity demand scalable, language-specific solutions. Finally, ethical concerns, such as bias in AI models and the loss of human judgment, remain underexplored in automated quality assessment.

This thesis aims to address these gaps by developing a multi-class classification model tailored for the Czech language media environment. By integrating textual analysis, metadata, and Czech-specific linguistic features, this work builds on existing research while tackling its limitations, particularly for news aggregation platforms.

■ 1.5 Contributions

Summarize your specific contributions and innovations Highlight what makes your approach unique

■ 2 How to write thesis in LaTeX

■ 2.1 Versioning with git

Write the LaTeX in such a way that it could be versioned by git, which will help when collaborating with other people. This means writing **one sentence per line**. Even when you use third-party platforms, such as the OverLeaf, you can still share the repository through Git.

■ 2.2 Forming paragraphs

A paragraph is formed in LaTeX by an uninterrupted block of non-empty lines. It is recommended to keep a single sentence per line (helps with versioning using git). A new paragraph is started after an empty line.

This is a new paragraph. It is strongly recommended to **avoid** the use of the *newline* (`\\`) feature of LaTeX for forming paragraphs as it doesn't format the new paragraph properly (no space at beginning of the new paragraph).

■ 2.3 Linguistic anti patterns

■ Narrative

We recommend to write your thesis in plural form of the first-person narrative in combination with passive tense, e.g.:

- We discourage the use of any other form, and/or
- any other form is discouraged, but **not**
- I discourage you from using the first-person narrative.

Moreover, avoid “instructional” or “teacher”-like style of writing, such as “**Now, we multiply the matrix $\mathbf{A}$ by the scalar c to get the scaled matrix $\mathbf{B}$.**” A better way of writing the same information would be e.g. “Now, the scaled matrix $\mathbf{B}$ is obtained by multiplying the matrix $\mathbf{A}$ by the scalar c .”

■ Pronouns

The use of pronouns (it, this, they) is strongly **discouraged**. Although, pronouns make it easier for you as a writer to form the flow of the text, pronouns also make it much more difficult for the reader to follow the text. The reader is forced to retain more of the context to substitute and understand what the author meant. Moreover, pronouns can easily become vague (there is more than one way how to interpret them) and can become invalid while making editorial changes to the text, i.e., when moving sentences around. A technical text should be written in a way that makes it as easy to read and comprehend as possible and as hard to misunderstand or misinterpret as possible at the same time.

■ 2.4 Mathematical notation with LaTeX

Take care to use the correct mathematical symbols and common ways of denoting mathematical concepts. Use bold fonts to visually distinguish vectors and matrices ($\mathbf{x}$, $\mathbf{A}$) and

scalars (k , N).

■ Common errors

A frequent error, carried over from programming languages, is using the asterisk symbol ($*$) to denote multiplication. The asterisk correctly denotes convolution. Similarly, the cross sign ($\times$) typically denotes the cross product (it can also be used for stating dimensions, such as $10\text{ m} \times 10\text{ m}$) and thus should not be used for scalar multiplication. In English mathematical notation, **scalar multiplication is typically not denoted at all**.

This custom may sometimes make it unclear whether a sequence of letters denotes multiplication of several scalars or a multi-letter variable, such as

$$T = T0 + coef f meas, \quad (2.1)$$

where the variables in this hypothetical equation are T , $T0$, $coef$ and $meas$. For this reason, **avoid using multi-letter variable naming** and strive to denote mathematical variables with single letters optionally with a lower or upper index, or other modifiers ($\hat{}$, $\bar{}$, etc.). The equation above could be modified to be

$$T = T_0 + cT_{\text{meas}}. \quad (2.2)$$

If the multiplication is still unclear (e.g. when multiplying many single-letter scalars), the $\cdot$ symbol may be used such as

$$P \cdot V = n \cdot R \cdot T. \quad (2.3)$$

■ Equations

Mathematical equations should be numbered and should be a part of a sentence. For example, a discrete LTI system update is described as

$$\mathbf{x}_{[k+1]} = \mathbf{A}\mathbf{x}_{[k]} + \mathbf{B}\mathbf{u}_{[k]}, \quad (2.4)$$

where $\mathbf{x}_{[k]} \in \mathbb{R}^m$ is the state vector at the sample k , $\mathbf{u}_{[k]} \in \mathbb{R}^n$ is the input vector, $\mathbf{A} \in \mathbb{R}^{m \times m}$ is the main system matrix, and $\mathbf{B} \in \mathbb{R}^{m \times n}$ is the system input matrix. Proper punctuation should be used after the equation, as if it were an ordinary object in the sentence.

Do not put any empty lines before the equation. If the sentence that the equation is a part of continues after the equation (as is the case here), do not put empty lines after the equation either. That would create a new paragraph mid-sentence. **For an example of how not to do it, the equation**

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (2.5)$$

describes the logistic function often used in machine learning. Observe how a new paragraph is created for the equation and then for this block of text (compare with the proper typesetting above). Not only does this not look correct, it may also cause incorrect page breaking.

■ 2.5 Using footnotes

Do not be afraid to use footnotes for additional information, such as http links¹. We use footnote links whenever we want to *point* to a website, rather than to cite it as a source. Like with everything, do not overdo it.

■ 2.6 Referencing document elements

LaTeX allows you to dynamically reference to parts of the documents, such as

- figures: Fig. 2.4, Figure 2.4,
- equations: eq. (2.4), (2.4),
- code: Lst. 2.1,
- and any other object that can contain a `\label`.

Check the section in the `document_setup.tex` that contains useful macros for unifying the references:

```
\newcommand{\reffig}[1]{Fig.~\ref{#1}}
\newcommand{\reflst}[1]{Lst.~\ref{#1}}
\newcommand{\refalg}[1]{Alg.~\ref{#1}}
\newcommand{\refsec}[1]{Sec.~\ref{#1}}
\newcommand{\reftab}[1]{Table~\ref{#1}}
\newcommand{\refeq}[1]{\eqref{#1}}
```

Listing 2.1: LaTeX macros for referencing to document elements.

■ 2.7 Abbreviations with Acronym

Abbreviations are handled by the *acronym* package. Example sentence with abbreviations: “Unmanned Aerial Vehicle (UAV) is a flying vehicle that commonly uses Light Detection and Ranging (LiDAR) and Global Positioning System (GPS) receiver”. Note that the acronyms are only explained once in the document by default. It is good practice to re-explain acronyms used both in the abstract and the rest of the document as the abstract is often presented separately. This can be achieved by resetting the internal status of the acronyms (“forgetting” that they were explained) using the `\acresetall` command after the abstract. Please, read the documentation².

■ 2.8 Units of measurements with Siunitx

Typesetting of units has never been more accessible with the Siunitx package. Acceleration is measured in ms^{-2} . Gravity accelerates objects at a rate $\approx 9.81 \text{ ms}^{-2}$ near the sea level. You can define your units if you want.

■ 2.9 Hyphens and dashes

Hyphens and dashes are the various form of the symbol “-” used in many situations. There are also various ways how to typeset the symbol in LaTeX.

- The *hyphen* is used to compound words, e.g., “the eye-opener”. The hyphen is typeset as a single *minus/hyphen* character: -.

¹This repository: <https://github.com/ctu-mrs/thesis-template>.

²Acronym package: <http://mirrors.ctan.org/macros/latex/contrib/acronym/acronym.pdf>

- The *en-dash* is used to specify ranges of values, e.g., “between 2–10”. The en-dash is typeset as two consecutive hyphens characters: --.
- The *em-dash* is used to separate complex sentences in place of commas, parenthesis and colons — each with its particular rules. The em-dash is typeset as three consecutive hyphens characters: ---.

Check the <https://www.thepunctuationguide.com/> for all the details.

■ 2.10 Double quotation marks

“Double quotes” in English are composed of a pair of opening (“) and closing (”) symbols. The opening symbol is typeset as two backtick characters: ‘‘ (typically below the Esc key on the English keyboard), and the closing quotes as two apostrophes: ’’. The LaTeX engine will convert them automatically to the opening and closing symbols. A more robust solution is to use the `csquotes` package and the `\enquote` command which also takes care of nested quoting and other peculiarities.

■ 2.11 2D Diagrams with Tikz

Tikz is a powerful tool for drawing 2D (and 3D) shapes and diagrams. Check the documentation and examples: https://www.overleaf.com/learn/latex/TikZ_package. The benefit of using *Tikz*, instead of some other third-party drawing program, are:

- fonts are the same as in LaTeX,
- you can typeset math in LaTeX,
- you can use references to other parts of your document,
- you can version the image in git,
- the images are easily adjustable while editing your document.

Check Fig. 2.1 for example.




Figure 2.1: Example of a 2D diagram using tikz *PGFPlots*.

■ 2.12 Data plots with PGFPlots

PGFPlots produces nice 2D and 3D data plots from data stored in CSV. The plot parameters can be versioned and easily adjusted by editing the plot definition file.

- Documentation and manual: <https://ctan.org/pkg/pgfplots>
- Compile the plots individually and then include the pdfs because it can take longer.
- Example located in `fig/plots/example_plot`, see Fig. 2.2.

- You could include the latex file directly. However, it will take longer to compile, and platforms such as Overleaf can have a problem with that.




Figure 2.2: Example of a 2D plot using *PGFPlots*.

■ 2.13 3D Plots with Sketch

Sketch is a tool for defining a 3D scene using simple descriptive language. The 3D scene is then converted to *Tikz*, which is later compiled to pdf. The benefits of using *Sketch* are similar to using *Tikz*: LaTeX fonts, versioning using git, and cleanness of the result. See the example image in Fig. 2.3.

- Documentation and manual: <http://www.frontiernet.net/~eugene.ressler/>
- Cross-compilation from *Sketch* to *pdf* using the `fig/sketch/compile_sketch.sh` script.




Figure 2.3: Depiction of the used coordinate systems. The image was drawn using *Sketch*.

■ 2.14 Image collages with Subfig

We recommend using the `subfig` package, which provides the `\subfloat` command. It is more versatile than the simpler `subcaption` package. Check Fig. 2.4 for an example.

■ 2.15 Citations with Biblatex

Biblatex is probably the most powerful citation package for LaTeX. It consumes the standard `.bib` file. However, it can sort and filter the citations using the `keywords` tag. Citing references is done using the `cite` command, e.g., **baca2021mrs**. You can also define some nice citation boxes, such as this one:

baca2021mrs

■ 2.16 Image overlays with Tikz

Tikz is very useful to create custom image overlays. The overlay can be set such that the image is spanned by Cartesian coordinates $(x, y) \in [0, 1]^2$. Example can be seen in Fig. 2.5.




Figure 2.4: The caption should mention both subfigures, the Fig. 2.4a and the Fig. 2.4b. You can just refer to them as (a) and (b) in the main Figure’s caption, but beware, you need to keep it correct as you edit.




Figure 2.5: Example of using Tikz for image overlays. (a) shows a final product, (b) shows a grid useful for nailing down the coordinates.

■ 2.17 General tips

In general, strive to make the paper easy to read and understand, and hard to misunderstand or misinterpret. Here are some more specific tips on how to achieve that (and other general suggestions).

- **Be consistent.** This applies in all contexts. For example, if you decide to use the name “LiDAR”, do not mix it with “LIDAR” or “Lidar”, do not mix different mathematical notations, ensure your Figures have the same style and use the same graphics for the same concepts, etc.
- After you finish writing or modifying any of:
 - a sentence,
 - a paragraph,
 - a section/chapter,
 - the whole paper/thesis,**re-read it** to make sure that it makes sense, it is coherent and correct, and doesn’t contain typos.
- If you’re using a LLM-based tool (ChatGPT etc.) for grammar-proofing or even formulation of sentences, **do not just copy-paste its response** to your query. The previous rule applies doubly here. LLMs tend to often produce confident-sounding nonsense, sentences with reformulated duplicated content, or with a slightly changed meaning. They

are a good tool to get inspiration to start writing about a subject, for grammar-checking, or for finding alternative, nice-sounding formulations, but they can lie or warp facts — take care when using them!

■ 3 Conclusion

Summarize the achieved results. Can be similar as an abstract or an introduction, however, it should be written in past tense.

■ 4 References

■ A Appendix A